

## What is TensorFlow?

- a) Library ✗
- b) Package ✗
- c) Ecosystem ✗
- d) Framework ➔

## What is a framework?

A framework in software (including DL) is a comprehensive set of tools, libraries, rules and structures that provides everything you need to build, train and deploy models or applications

TENSORBOARD → Track models' performance leveraging Tensorboard.

Some framework examples: TensorFlow, PyTorch, Django (Web), React.js,

Spring (Java)

is a deep learning framework.

- TensorFlow (TF) is an open-source developed by Google Brain in 2015
  - [for their internal research & production]
- TF is widely used for building DL models due to its flexibility, scalability and ability to run on CPUs: central processing units  
GPUs: Graphics Processing Units  
TPUs: Tensor Processing Units (by Google)
- TF supports building models at multiple abstraction levels
  - from low level operations to high level APIs like tf.Keras



## TENSORFLOW (low level)

```
python
import tensorflow as tf
import numpy as np
# Dummy data
X = np.random.randn(100, 2).astype(np.float32)
y = (np.sum(X, axis=1) > 0).astype(np.float32).reshape(-1, 1)
```

```
python
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
```

on 26 Oct

1st Nov.

## TF-KERAS HIGH LEVEL API

```
import numpy as np

# Dummy data
X = np.random.randn(100, 2).astype(np.float32)
y = (np.sum(X, axis=1) > 0).astype(np.float32).reshape(-1, 1)

# Initialize weights and biases
W1 = tf.Variable(tf.random.normal([2, 4]))
b1 = tf.Variable(tf.zeros([4]))
W2 = tf.Variable(tf.random.normal([4, 1]))
b2 = tf.Variable(tf.zeros([1]))

# Define model
def forward_pass(X):
    z1 = tf.matmul(X, W1) + b1
    a1 = tf.nn.relu(z1)
    z2 = tf.matmul(a1, W2) + b2
    output = tf.nn.sigmoid(z2)
    return output

# Loss and optimizer
def loss_fn(y_true, y_pred):
    return tf.reduce_mean(tf.keras.losses.binary_crossentropy(y_true, y_pred))

optimizer = tf.optimizers.Adam(learning_rate=0.01)

# Training Loop
for epoch in range(100):
    with tf.GradientTape() as tape:
        y_pred = forward_pass(X)
        loss = loss_fn(y, y_pred)
    grads = tape.gradient(loss, [W1, b1, W2, b2])
    optimizer.apply_gradients(zip(grads, [W1, b1, W2, b2]))
    if epoch % 10 == 0:
        print(f"Epoch {epoch}, Loss: {loss.numpy():.4f}")

from tensorflow.keras import layers
import numpy as np

# Dummy data
X = np.random.randn(100, 2)
y = (np.sum(X, axis=1) > 0).astype(np.float32)

# Build model
model = keras.Sequential([
    layers.Dense(4, activation='relu', input_shape=(2,)),
    layers.Dense(1, activation='sigmoid')
])

# Compile
model.compile(optimizer='adam',
               loss='binary_crossentropy',
               metrics=['accuracy'])

# Train
model.fit(X, y, epochs=100, verbose=1)
```

## MLP from scratch - APC Model (even lower level)



250ch

```

INPUT LAYER ---> (1) HIDDEN LAYER ---> OUTPUT LAYER

class NeuralNetwork:
    def __init__(self, input_size, hidden_size, output_size, learning_rate = 0.01, epochs=1000):
        self.input_size = input_size # no. of input vars or Features in the Input Layer
        self.hidden_size = hidden_size # no. of neurons in the Hidden Layer (In here we have only one Hidden Layer)
        self.output_size = output_size # no. of neurons in the output layer
        self.learning_rate = learning_rate # learning rate for the gradient descent step
        self.epochs = epochs # no. of training epochs

    ##### INITIALIZE WEIGHTS AND BIASES #####
    #### INPUT LAYER TO HIDDEN LAYER #####
    #####
    self.W1 = np.random.rand(self.input_size, self.hidden_size)*0.01 # random weights drawn from std. normal distribution
    self.b1 = np.zeros((1, self.hidden_size)) # zero bias values for the neurons in the hidden layer

    #####
    #### HIDDEN LAYER TO OUTPUT LAYER #####
    #####
    self.W2 = np.random.rand(self.hidden_size, self.output_size)*0.01 # random weights drawn from std. normal distribution
    self.b2 = np.zeros((1, self.output_size)) # zero bias values for the neurons in the output layer

    #####
    #### ACCURACY HISTORY FOR PREDICTION #####
    self.loss_history = [] # Empty list initialized to store the losses during training epochs
    self.accuracy_history = [] # Empty list initialized to store the accuracy values during training epochs

    #####
    #### ADD SOME ACTIVATION FUNCTIONS #####
    #####
    # ReLU activation function
    def relu(self, x):
        return np.maximum(0,x)

    # Derivative of ReLU for backpropagation
    def relu_derivative(self, x):
        return np.where(x>0, 1, 0)

    # Softmax activation function
    def softmax(self, x):
        exp_values = np.exp(x - np.mean(x, axis=1, keepdims=True)) # subtract max for numerical stability
        return exp_values/np.sum(exp_values, axis=1, keepdims=True)

    #####
    # 1. FORWARD PROPAGATION #####
    #####
    def forward(self,X):
        #####
        # INPUT LAYER TO HIDDEN LAYER #####
        #####
        self.z1 = np.dot(X, self.W1) + self.b1 # compute z = w1X + b1
        self.a1 = self.softmax(self.z1) # pass z1 to ReLU activation function to get output as a1
        return a1

        #####
        # HIDDEN LAYER TO OUTPUT LAYER #####
        #####
        self.z2 = np.dot(self.a1, self.W2) + self.b2 # compute z = w2a1 + b2
        self.a2 = self.softmax(self.z2) # pass z2 to Softmax activation function to get output as a2
        return a2, self.a1

    #####
    # 2. COMPUTE LOSS & ACCURACY #####
    #####
    # Cross-entropy loss for multi-class classification
    def compute_loss(self, y_true, prob):
        loss = -np.mean(np.multiply(y_true * np.log(prob), axis=1))
        return loss

    # Compute accuracy
    def compute_accuracy(self, y_true, prob):
        predictions = np.argmax(prob, axis=1)
        predictions = np.argmax(prob, axis=1)
        true_labels = np.argmax(y_true, axis=1)
        true_labels = np.argmax(y_true, axis=1)
        return np.mean(predictions == true_labels)

    #####
    # 3. BACKPROPAGATION #####
    #####
    def backward(self, X, y):

        # Using Batch Gradient Descent (BGD)
        # Number of rows/training examples
        n = X.shape[0] # all rows in the data

        #####
        # Gradients of the loss w.r.t. weights and biases of the output layer #####
        #####
        # delta_t = np.sum(delta_t, 0) # Average of the output layer
        # delta_t = np.dot(delta_t, 1/n) # Gradient of the loss w.r.t. weights of the output layer --> did
        # did = np.sum(delta_t, 0, keepdims=True) # Gradient of the loss w.r.t. bias of the output layer --> did

        #####
        # Gradients of the loss w.r.t. weights and biases of the hidden layer #####
        #####
        # delta_t = np.sum(delta_t, 0) # Average of the hidden layer
        # delta_t = np.dot(delta_t, 1/n) # Gradient of the loss w.r.t. weights of the hidden layer --> did
        # did = np.sum(delta_t, 0, keepdims=True) # Gradient of the loss w.r.t. bias of the hidden layer --> did

        #####
        # Updating the parameters #####
        # -> weights and biases across the layers (Hidden layer & Output layer)
        #####
        self.W2 -= self.learning_rate * did
        self.b2 -= self.learning_rate * did
        self.W1 -= self.learning_rate * did
        self.b1 -= self.learning_rate * did
        self.a1 -= self.learning_rate * did

```

```

dD = np.sum(dDelta, axis=0, keepdims=True) # gradient of the loss w.r.t. bias of the hidden layer --> dB2

##### Updating the parameters --> weights and biases across the layers (hidden layer & output layer)
self.w2 -= self.learning_rate * dB2
self.b2 -= self.learning_rate * bD2
self.w1 -= self.learning_rate * dB1
self.b1 -= self.learning_rate * bD1

#####
# MLP TRAINING
#####
def train(self, X,y):
    for epoch in range(self.epochs):

        ##### 1. Forward pass
        pred = self.forward() #computes the prediction using forward pass method

        ##### 2. Compute the loss and accuracy
        loss = self.computeLoss(y, pred)
        accuracy = self.computeAccuracy(y, pred)
        self.lossHistory.append(loss)
        self.accuracyHistory.append(accuracy)

        ##### 3. Backward Pass
        self.backward(X,y)

        ##### Print the progress after every 10 epochs
        if epoch % 10 == 0:
            print(f"Training metrics: Epoch: {epoch}, Loss: {loss:.4f}, Accuracy: {accuracy:.4f}")


```

Below is the ML code snippet for this page

<https://www.tensorflow.org/>

The screenshot shows the TensorFlow website's main landing page. At the top, there's a navigation bar with links for 'Install', 'Learn', 'API', 'Ecosystem', and 'More'. Below the navigation is a search bar and a GitHub link. The main content area features a large heading 'Get started with TensorFlow'. Below the heading, a brief description states: 'TensorFlow makes it easy to create ML models that can run in any environment. Learn how to use the intuitive APIs through interactive code samples.' To the left of the text is a small icon of a person looking at a computer screen. On the right, there's a code editor window displaying Python code for a neural network to classify MNIST digits. A 'Run quickstart' button is at the bottom of the code window. The background has a light gray gradient with some abstract white and orange lines.

<https://github.com/tensorflow/tensorflow/tree/master/tensorflow>

The screenshot shows a terminal window running on a Mac OS X desktop. It displays the Python code for a neural network to classify MNIST digits. The code is identical to the one shown on the TensorFlow website. A handwritten note '(15)' is circled in blue ink next to the line `model.compile(optimizer='adam',` and another note 'lines of code' is written below the code block. At the bottom of the terminal window is a 'Run quickstart' button.

```

import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

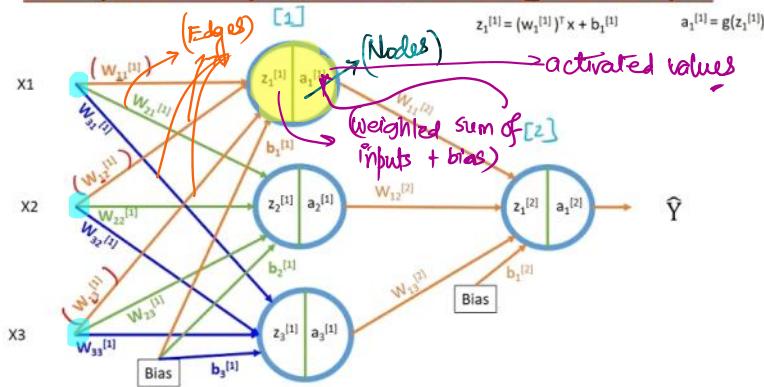
model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)

```

Tensor Flow

What do you mean by 'Tensor' and 'Flow' in TensorFlow?

## Compute the prediction for single example



In any n/w graph, there are nodes and edges:

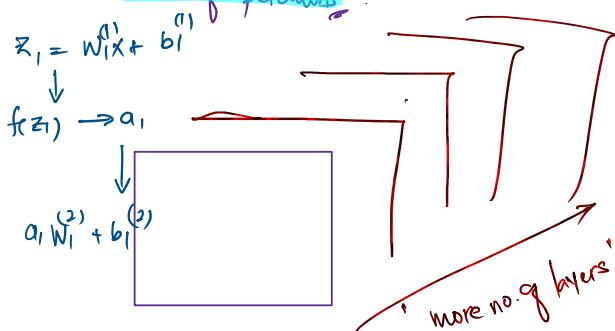
$$W = \begin{bmatrix} W_{11} & W_{21} & W_{31} & \dots \\ W_{12} & W_{22} & W_{32} & \dots \\ \vdots & & & \\ W_{15} & W_{25} & W_{35} & \dots \end{bmatrix}$$

represented by an array.

- edges in any n/w graph represent  $\rightarrow$  multi-dimensional arrays  
 $\downarrow$   
 called as TENSOR.

Tensor: represents multi-dimensional data structure (ndim arrays)

Flow: represents the way data moves through the n/w due to series of operations



## # KERAS

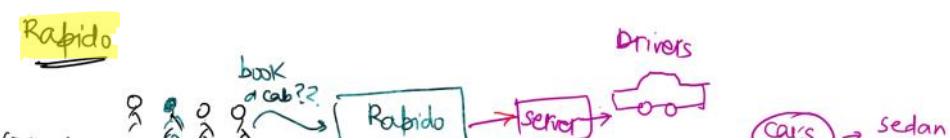
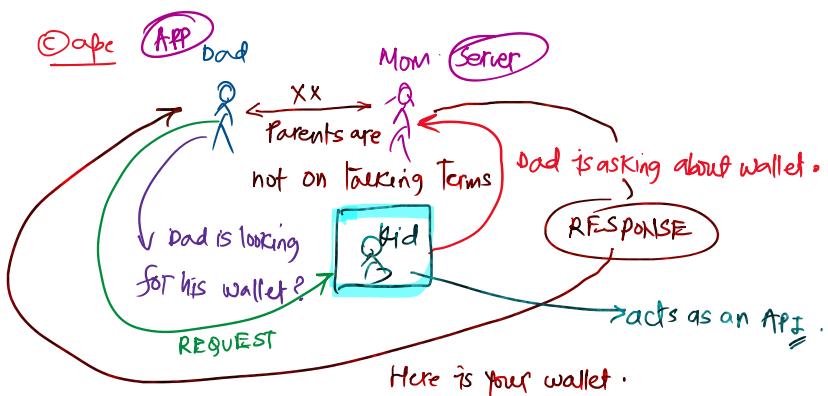
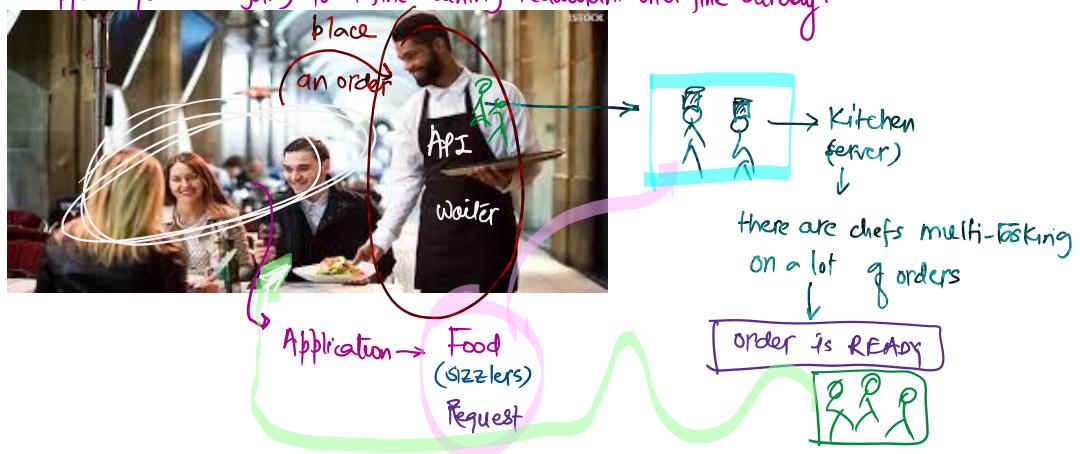
What is Keras?

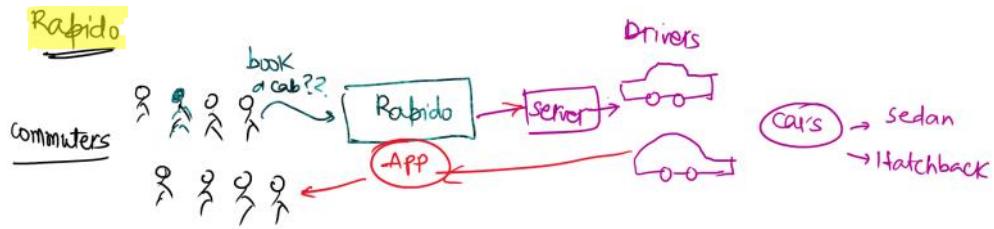
- library / API layer  
 (High level neural n/w API)
- It is an open source high level neural n/w API (Application Programming Interface) developed by Francois chollet and it released in 2015.

## What's an API ??

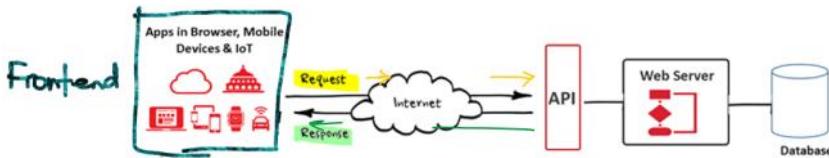
- stands for Application Programming Interface.

Suppose you are going to a fine-dining restaurant on a fine Sunday.





An API is a software which can be used by other softwares to communicate with other softwares or even hardwares. It acts as a bridge between different softwares and devices.



API is a software intermediary that allows two or more applications to talk to each other.

### Why do we need Keras?

① User-friendly: Easy to use and understand → making it ideal platform for the beginners.

② Modularity: Building blocks for neural net models that's been saved after previously being trained on a large-scale datasets.

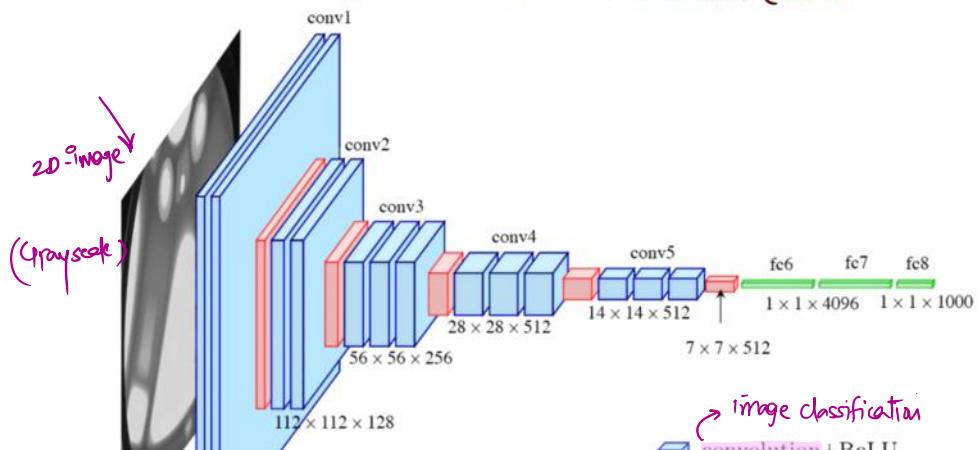
**ChatGPT** → Can run on android/iOS phones/tablets  
Can run on very low configured laptops/computers

Note: Keras does provide several pre-trained models

- ResNet 50
- YOLO v12
- DeepSeek FF
- VGG 16
- YOLO v08

# VGG16 is a deep convolutional neural network (CNN)

# VGG16 is a deep convolutional neural network (CNN)





Manav - difference between API and ABI

### TENSORFLOW vs KERAS

#### TENSORFLOW

- Type: Deep learning framework

- Level: Low-level + High-level (via Keras)

- Origin: Developed by Google Brain in 2015

- Purpose: Build, train, deploy → DL models.  
(R&D) work → Tensorflow

- Flexibility: very flexible, allows for custom operations and full control

- Standalone: Yes, can be run without Keras.

#### KERAS

- High level API for DL  
or  
High level library

- High-level only.

- Developed by Francois Chollet in 2015 too.

- Quickly build & train  
or for experimenting  
- Quick & dirty POC

- Designed for simplicity,  
and limited or less  
control over low-level operations

- Keras is tightly integrated  
with Tensor Flow

-MNIST Use-case - we'll build using TF first only & then using TF-Keras